

Grafana Labs

Products

Open source

Solutions

Learn

Docs

Company

🔍

Downloads

Contact us

Sign in

Grafana Cloud

Grafana

Grafana Loki

Grafana Mimir

Grafana OnCall

Grafana Tempo

Grafana k6

More docs ▾

Grafana Labs documentation

🔍 Search docs

Product

Grafana k6 ▾

Viewing: v0.51.x (latest)

Find another version

Grafana k6 documentation

> Get started

> Set up

> Using k6

HTTP Requests

> Metrics

Checks

Thresholds

> Options

Test lifecycle

Documentation > Grafana k6 > Using k6 > Scenarios > Executors > Constant arrival rate

Open source

## Constant arrival rate

With the `constant-arrival-rate` executor, k6 starts a fixed number of iterations over a specified period of time. It is an open-model executor, meaning iterations start independently of system response (for details, read [Open and Closed models](#)).

This executor continues to start iterations at the given rate as long as VUs are available. The time to execute an iteration can vary with test logic or the system-under-test response time. To compensate for this, the executor starts a varied number of VUs to meet the configured iteration rate. For explanations of how allocation works, read [Arrival-rate VU allocation](#).

NOTE

Iteration starts are spaced fractionally. Iterations do not start at exactly the same time. At a rate of 10 with a `timeUnit` of 1s, each iteration starts about every tenth of a second (that is, each 100ms).

## Options

Besides the [common configuration options](#), this executor has the following options:

Option	Type	Description	Default
<code>duration</code> <sup>(required)</sup>	string	Total scenario duration (excluding <code>gracefulStop</code> ).	-
<code>rate</code> <sup>(required)</sup>	integer	Number of iterations to start during each <code>timeUnit</code> period.	-
<code>preAllocatedVUs</code> <sup>(required)</sup>	integer	Number of VUs to pre-allocate before test start to preserve runtime resources.	-
<code>timeUnit</code>	string	Period of time to apply the <code>rate</code> value.	"1s"
<code>maxVUs</code>	integer	Maximum number of VUs to allow during the test run.	If unset, same as <code>preAllocatedVUs</code>

## When to use

When you want iterations to remain constant, independent of the performance of the system under test. This approach is useful for a more accurate representation of RPS, for example.

NOTE

**Don't put sleep at the end of an iteration.**

The arrival-rate executors already pace the iteration rate through the `rate` and `timeUnit` properties. So it's unnecessary to use a `sleep()` function at the end of the VU code.

## Example

This example schedules a constant rate of 30 iterations per second for 30 seconds. It pre-allocates 2 VUs, and allows k6 to dynamically schedule up to 50 VUs as needed.

JavaScript

Copy

```
import http from 'k6/http';

export const options = {
  discardResponseBodies: true,
  scenarios: {
    contacts: {
      executor: 'constant-arrival-rate',

      // How long the test lasts
      duration: '30s',

      // How many iterations per timeUnit
      rate: 30,
```

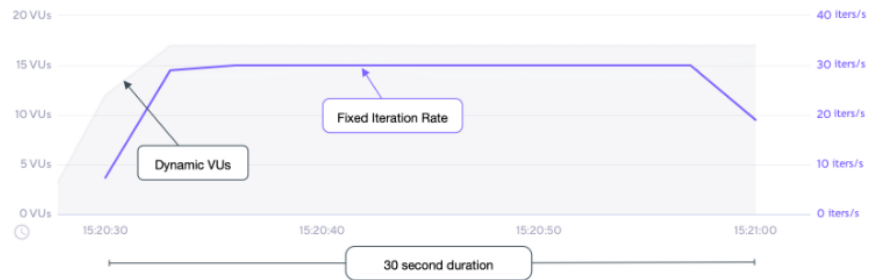
```
// Start 'rate' iterations per second
timeUnit: '1s',

// Pre-allocate 2 VUs before starting the test
preAllocatedVUs: 2,
```

[Expand code](#)

## Observations

The following graph depicts the performance of the `example` script:



Based upon our test scenario inputs and results:

- The desired rate of 30 iterations started every 1 second is achieved and maintained for the majority of the test.
- The test scenario runs for the specified 30 second duration.
- Having started with 2 VUs (as specified by the `preAllocatedVUs` option), k6 automatically adjusts the number of VUs to achieve the desired rate, up to the `maxVUs`. For this test, this ended up as 17 VUs.
- The number of VUs to achieve the desired rate varies depending on how long each iteration takes to execute. For this test definition, if it would take exactly 1 second, then 30 VUs would be needed. However, as it takes less than 1 second, then less VUs are needed.

Using too low of a `preAllocatedVUs` setting will reduce the test duration at the desired rate, as resources need to continually be allocated to achieve the rate.

## Was this page helpful?

[Suggest an edit](#)[Contribute to docs](#)[Report a problem](#)

## Related documentation

[Community](#)[Support](#)[API load testing](#)[Performance testing with Grafana Cloud k6](#)[Use the test builder](#)

## Related resources from Grafana Labs


Additional helpful documentation, links, and articles:

Video



Performance testing and observability in Grafana Cloud

In this webinar, learn how Grafana Cloud k6 offers you the best developer experience for performance testing.



User-centered observability: load testing, real user monitoring, and synthetics

Learn how to use load testing, synthetic monitoring, and real user monitoring (RUM) to understand end users' experience of your apps. Watch on demand.



Grafana

- Overview
- Deployment options
- Plugins
- Dashboards

Products

- Grafana Cloud
  - Grafana Cloud Status
- Grafana Enterprise Stack
- Grafana Cloud Application Observability
- Grafana Cloud Frontend Observability
- Grafana Cloud IRM
- Grafana Cloud k6
- Grafana Cloud Logs
- Grafana Cloud Metrics
- Grafana Cloud Profiles
- Grafana Cloud Synthetic Monitoring
- Grafana SLO

Open Source

- Grafana
- Grafana Loki
- Grafana Mimir
- Grafana OnCall
- Grafana Tempo
- Grafana Agent
- Grafana Alloy
- Grafana k6
- Prometheus
- Grafana Faro
- Grafana Pyroscope
- Grafana Beyla
- OpenTelemetry
- Grafana Tanka
- Graphite
- [GitHub](#)

Learn

- Grafana Labs blog
- Documentation
- Downloads
- Community
- Community forums
- Community Slack
- Grafana Champions
- Community organizers
- Grafana ObservabilityCON
- GrafanaCON 2024
- The Golden Grot Awards
- Successes
- Workshops
- Videos
- OSS vs Cloud
- Load testing

Company

- The team
- Press
- Careers
- Events
- Partnerships
- Contact
- Getting help
- Merch